

**ŽILINSKÁ UNIVERZITA V ŽILINE**  
**FAKULTA RIADENIA A INFORMATIKY**

**Projektovanie sietí 1**  
**BGP**

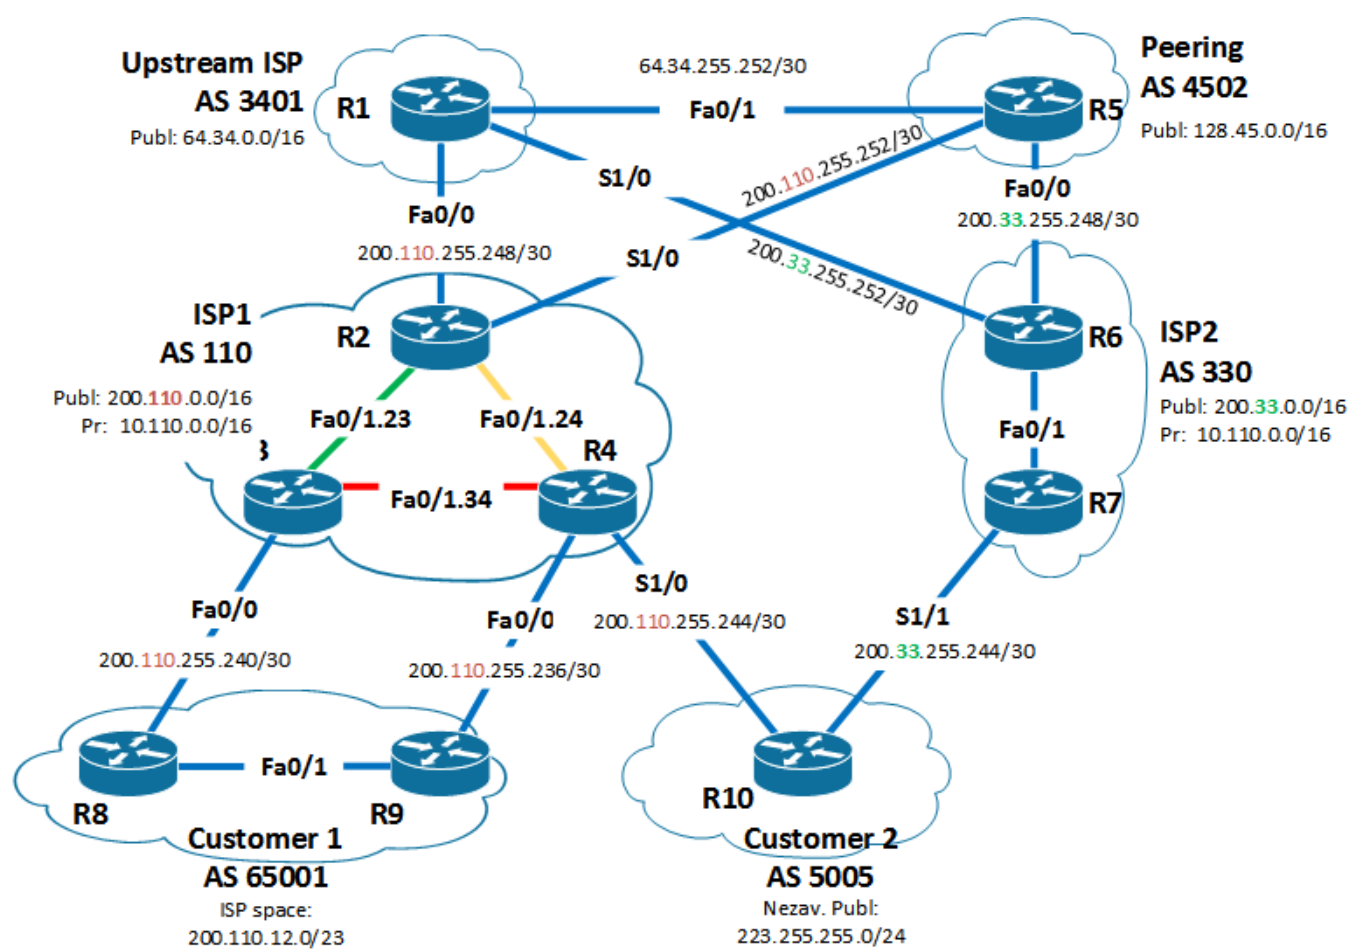
## Obsah

|   |    |
|---|----|
| 1. Zadanie .....  | 3  |
| 2. Fyzická topológia .....  | 3  |
| 3. Adresný plán .....   | 4  |
| 4. ISP konektivita a peering .....  | 5  |
| 4.1 Použitie IS-IS pre vnútorné smerovanie v rámci AS.....                            | 5  |
| 4.2 Distribúcia smerovacích záznamov z každého AS.....                                | 5  |
| 4.3 Zabezpečiť, aby interné ISP adresy neboli propagované.....                        | 7  |
| 4.4 Odstránenie privátnych AS .....   | 7  |
| 4.5 Sumarizácia záznamov .....  | 8  |
| 4.6 Kontrola konektivity .....  | 9  |
| 5. ISP politika .....   | 10 |
| 5.1 Primárne linky R3 – R8 a R4 – R10.....  | 10 |
| 5.2 AS 5005 nesmie byť nikdy transit .....  | 13 |
| 5.3 Peering iba pre ISP1 a ISP2.....  | 14 |
| 5.4 Distribuovať iba default, AS 5005 a peering prefixy do AS 65001 .....             | 15 |
| 5.5 Overiť, či je možné odkloniť celú prevádzku na linke R4-R10 v prípade údržby..... | 16 |
| 5.6 Overiť funkčnosť nastavenia politiky vhodnými výpadkami liniek .....              | 18 |

## 1. Zadanie

Cieľom cvičenia bolo oboznámiť sa so smerovacím protokolom BGP a jeho využitím pre internet peering.

## 2. Fyzická topológia



### 3. Adresný plán

| ROUTER | INTERFACE | ADRESA          | MASKA           |
|--------|-----------|-----------------|-----------------|
| R1     | Fa0/0     | 200.110.255.249 | 255.255.255.252 |
|        | Fa0/1     | 64.34.255.253   | 255.255.255.252 |
|        | S1/0      | 200.33.255.253  | 255.255.255.252 |
|        | Loopback0 | 10.255.255.1    | 255.255.255.255 |
|        | Loopback1 | 64.34.0.1       | 255.255.255.0   |
| R2     | Fa0/0     | 200.110.255.250 | 255.255.255.252 |
|        | Fa0/1.23  | 10.110.23.2     | 255.255.255.0   |
|        | Fa0/1.24  | 10.110.24.2     | 255.255.255.0   |
|        | S1/0      | 200.110.255.253 | 255.255.255.252 |
|        | Loopback0 | 10.255.255.2    | 255.255.255.255 |
|        | Loopback1 | 200.110.2.1     | 255.255.255.0   |
| R3     | Fa0/0     | 200.110.255.241 | 255.255.255.252 |
|        | Fa0/1.23  | 10.110.23.3     | 255.255.255.0   |
|        | Fa0/1.34  | 10.110.34.3     | 255.255.255.0   |
|        | Loopback0 | 10.255.255.3    | 255.255.255.255 |
|        | Loopback1 | 200.110.3.1     | 255.255.255.0   |
| R4     | Fa0/0     | 200.110.255.237 | 255.255.255.252 |
|        | Fa0/1.24  | 10.110.24.4     | 255.255.255.0   |
|        | Fa0/1.34  | 10.110.34.4     | 255.255.255.0   |
|        | S1/0      | 200.110.255.245 | 255.255.255.252 |
|        | Loopback0 | 10.255.255.4    | 255.255.255.255 |
|        | Loopback1 | 200.110.4.1     | 255.255.255.0   |
| R5     | Fa0/0     | 200.33.255.249  | 255.255.255.252 |
|        | Fa0/1     | 64.34.255.254   | 255.255.255.252 |
|        | S1/0      | 200.110.255.254 | 255.255.255.252 |
|        | Loopback0 | 10.255.255.5    | 255.255.255.255 |
|        | Loopback1 | 128.45.0.1      | 255.255.255.0   |
| R6     | Fa0/0     | 200.33.255.250  | 255.255.255.252 |
|        | Fa0/1     | 10.110.67.6     | 255.255.255.0   |
|        | S1/0      | 200.33.255.254  | 255.255.255.252 |
|        | Loopback0 | 10.255.255.6    | 255.255.255.255 |
|        | Loopback1 | 200.33.6.1      | 255.255.255.0   |
| R7     | Fa0/1     | 10.110.67.7     | 255.255.255.0   |
|        | S1/1      | 200.33.255.245  | 255.255.255.252 |
|        | Loopback0 | 10.255.255.7    | 255.255.255.255 |
|        | Loopback1 | 200.33.7.1      | 255.255.255.0   |
| R8     | Fa0/0     | 200.110.255.242 | 255.255.255.252 |
|        | Fa0/1     | 10.110.89.8     | 255.255.255.0   |
|        | Loopback0 | 10.255.255.8    | 255.255.255.255 |
|        | Loopback1 | 200.110.12.1    | 255.255.255.0   |
| R9     | Fa0/0     | 200.110.255.238 | 255.255.255.252 |
|        | Fa0/1     | 10.110.89.9     | 255.255.255.0   |
|        | Loopback0 | 10.255.255.9    | 255.255.255.255 |
|        | Loopback1 | 200.110.12.129  | 255.255.255.0   |
| R10    | S1/0      | 200.110.255.246 | 255.255.255.252 |
|        | S1/1      | 200.33.255.246  | 255.255.255.252 |
|        | Loopback0 | 10.255.255.10   | 255.255.255.255 |
|        | Loopback1 | 223.255.255.1   | 255.255.255.0   |

## 4. ISP konektivita a peering

### 4.1 Použitie IS-IS pre vnútorné smerovanie v rámci AS

Na zabezpečenie vnútornej konektivity v rámci AS 110, 330 a 65001 sme spustili protokol IS-IS. Podrobnejšiu konfiguráciu a nastavenia neuvádzame, nakoľko sme sa tomuto protokolu venovali na predošlých dvoch cvičeniach.

P2P prepojenia v ISP1 na smerovači R2

```
3R2#sh run | sec interface FastEthernet0/1.23
```

```
interface FastEthernet0/1.23  
encapsulation dot1Q 23  
ip address 10.110.23.2 255.255.255.0  
ip router isis  
isis network point-to-point
```

```
3R2#sh run | sec interface FastEthernet0/1.24
```

```
interface FastEthernet0/1.24  
encapsulation dot1Q 24  
ip address 10.110.24.2 255.255.255.0  
ip router isis  
isis network point-to-point
```

P2P prepojenia v ISP2 na smerovači R6

```
3R6#sh run | sec interface FastEthernet0/1
```

```
interface FastEthernet0/1  
ip address 10.110.67.6 255.255.255.0  
ip router isis  
isis network point-to-point
```

### 4.2 Distribúcia smerovacích záznamov z každého AS

Aby sme mohli začať distribuovať záznamy, bolo potrebné najskôr spustiť smerovací protokol BGP na každom smerovači.

```
router bgp #číslo_AS
```

Po spustení smerovacieho protokolu bolo potrebné nadviazať susedské vzťahy, aby bola zabezpečená konektivita podľa zadania. Smerovače, ktoré nadväzovali susedstvá so smerovačmi v inom AS využívali externé BGP (eBGP), pričom smerovače v rámci jedného AS využívali interné BGP (iBGP).

Príklad konfigurácie eBGP medzi R1 a R5

```
3R1: neighbor 64.34.255.254 remote-as 4502
```

```
3R5: neighbor 64.34.255.253 remote-as 3401
```

Pri konfigurácii eBGP využívame priame spojenie so susedom a do príkazu *neighbor* zadávame IP adresu jeho rozhrania. Konfigurácia iBGP je mierne odlišná, keďže v rámci AS nám beží IGP protokol, v našom prípade IS-IS. Aby sme zabezpečili plnú konektivitu aj v prípade výpadku niektorej linky vo vnútri AS, BGP susedské vzťahy nadväzujeme cez Loopbacky.

#### Príklad konfigurácie iBGP v rámci AS 110 (smerovač R2)

```
neighbor 10.255.255.3 remote-as 110
neighbor 10.255.255.3 update-source Loopback0
neighbor 10.255.255.4 remote-as 110
neighbor 10.255.255.4 update-source Loopback0
```

Posledným krokom je zapnutie ohlasovania požadovaných sietí. V našom prípade sme na každom smerovači prideliť IP adresu na Loopback1 a následne sme začali túto sieť oznamovať cez BGP pomocou príkazu *network*.

```
3R1: network 64.34.0.0 mask 255.255.255.0
3R2: network 200.110.2.0 mask 255.255.255.0
3R3: network 200.110.3.0 mask 255.255.255.0
3R4: network 200.110.4.0 mask 255.255.255.0
3R5: network 128.45.0.0 mask 255.255.255.0
3R6: network 200.33.6.0 mask 255.255.255.0
3R7: network 200.33.7.0 mask 255.255.255.0
3R8: network 200.110.12.1 mask 255.255.255.0
3R9: network 200.110.12.129 mask 255.255.255.0
3R10: network 223.255.255.0 mask 255.255.255.0
```

#### Kontrola konektivity a správnej konfigurácie pomocou *tc/sh* skriptu:

```
3R10(tcl)#foreach address {
+>64.34.0.1
+>200.110.2.1
+>200.110.3.1
+>200.110.4.1
+>128.45.0.1
+>200.33.6.1
+>200.33.7.1
+>200.110.12.1
+>200.110.12.129
+>} {
+>ping $address source 223.255.255.1 }
```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 64.34.0.1, timeout is 2 seconds:

Packet sent with a source address of 223.255.255.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 12/25/48 ms

Sending 5, 100-byte ICMP Echos to 200.110.2.1, timeout is 2 seconds:

Packet sent with a source address of 223.255.255.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/28/48 ms

```

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 200.110.3.1, timeout is 2 seconds:
Packet sent with a source address of 223.255.255.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/27/40 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 200.110.4.1, timeout is 2 seconds:
Packet sent with a source address of 223.255.255.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms
Sending 5, 100-byte ICMP Echos to 128.45.0.1, timeout is 2 seconds:
Packet sent with a source address of 223.255.255.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/40/68 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 200.33.6.1, timeout is 2 seconds:
Packet sent with a source address of 223.255.255.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/26/44 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 200.33.7.1, timeout is 2 seconds:
Packet sent with a source address of 223.255.255.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/6/28 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 200.110.12.1, timeout is 2 seconds:
Packet sent with a source address of 223.255.255.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 48/59/68 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 200.110.12.129, timeout is 2 seconds:
Packet sent with a source address of 223.255.255.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/24/36 ms

```

### 4.3 Zabezpečiť, aby interné ISP adresy neboli propagované

Pre vyriešenie tohto bodu zadania, sme na rozhraniach, ktoré smerovali von z AS ISP nenastavovali protokol IS-IS. Pre úplne zabezpečenie sme navyše nastavili tieto rozhrania ako pasívne.

### 4.4 Odstránenie privátnych AS

Keďže smerovače R8 a R9 sa nachádzali v privátnom AS 65001, bolo potrebné túto informáciu odstrániť, aby sa nešírila za hranice nášho ISP, ktorý sa nachádza v AS 110. Zmeny bolo potrebné vykonať na smerovačoch R2 (smerom na R1,R5) a R4 (smerom na R10).

```

3R2: neighbor 200.110.255.249 remove-private-as
      neighbor 200.110.255.254 remove-private-as
3R4: neighbor 200.110.255.246 remove-private-as

```

Overenie vykonáme pomocou príkazu *sh ip bgp* zo smerovača R5 smerom na Loopback1 smerovača R8, kde môžeme vidieť, že sa v zozname AS\_PATH nenachádza autonómny systém 65001.

**3R5#sh ip bgp 200.110.12.0**

BGP routing table entry for 200.110.12.0/25, version 14  
Paths: (2 available, best #2, table Default-IP-Routing-Table)

Advertised to update-groups:

1

**3401 110**

64.34.255.253 from 64.34.255.253 (64.34.0.1)

Origin IGP, localpref 100, valid, external

**110**

200.110.255.253 from 200.110.255.253 (200.110.2.1)

Origin IGP, localpref 100, valid, external, best

## 4.5 Sumarizácia záznamov

Pre sprehľadnenie smerovacích tabuliek BGP sme vykonali sumarizáciu záznamov. Verejné siete, ktoré ohlasujeme cez BGP, a ktoré sa nachádzajú vo vnútri AS 110 (smerovače R2-R4), sme sumarizovali na jednu sieť 200.110.0.0/21 a siete v AS 65001 sa sumarizovali na sieť 200.110.12.0/24.

Sumarizácia smerovačov R2-R4

aggregate-address 200.110.0.0 255.255.248.0 summary-only

Sumarizácia R8 a R9

aggregate-address 200.110.12.0 255.255.255.0 summary-only

Overenie vykonáme pomocou príkazu *sh ip bgp*:

**3R5#sh ip bgp**

| Network                 | Next Hop        | Metric | LocPrf | Weight | Path            |
|-------------------------|-----------------|--------|--------|--------|-----------------|
| * 64.34.0.0/24          | 200.33.255.250  |        |        | 0      | 330 3401 i      |
| *                       | 200.110.255.253 |        |        | 0      | 110 3401 i      |
| *>                      | 64.34.255.253   | 0      |        | 0      | 3401 i          |
| *> 128.45.0.0/24        | 0.0.0.0         | 0      |        | 32768  | i               |
| * 200.33.6.0            | 64.34.255.253   |        |        | 0      | 3401 330 i      |
| *>                      | 200.33.255.250  | 0      |        | 0      | 330 i           |
| * 200.33.7.0            | 64.34.255.253   |        |        | 0      | 3401 330 i      |
| *>                      | 200.33.255.250  |        |        | 0      | 330 i           |
| * <b>200.110.0.0/21</b> | 64.34.255.253   |        |        | 0      | 3401 110 i      |
| *>                      | 200.110.255.253 | 0      |        | 0      | 110 i           |
| *                       | 200.33.255.250  |        |        | 0      | 330 3401 110 i  |
| *> <b>200.110.12.0</b>  | 200.110.255.253 |        |        | 0      | 110 i           |
| *                       | 64.34.255.253   |        |        | 0      | 3401 110 i      |
| * 223.255.255.0         | 64.34.255.253   |        |        | 0      | 3401 110 5005 i |
| *>                      | 200.110.255.253 |        |        | 0      | 110 5005 i      |



## 4.6 Kontrola konektivity

```
3R5(tcl)#foreach address {
```

```
+>64.34.0.1
```

```
+>200.110.2.1
```

```
+>200.110.3.1
```

```
+>200.110.4.1
```

```
+>200.33.6.1
```

```
+>200.33.7.1
```

```
+>200.110.12.1
```

```
+>200.110.12.129
```

```
+>223.255.255.1
```

```
+>} {
```

```
+>ping $address source 128.45.0.1 }
```

Sending 5, 100-byte ICMP Echos to 64.34.0.1, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 20/30/44 ms

Sending 5, 100-byte ICMP Echos to 200.110.2.1, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms

Sending 5, 100-byte ICMP Echos to 200.110.3.1, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/30/48 ms

Sending 5, 100-byte ICMP Echos to 200.110.4.1, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 20/28/32 ms

Sending 5, 100-byte ICMP Echos to 200.33.6.1, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 16/32/44 ms

Sending 5, 100-byte ICMP Echos to 200.33.7.1, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 48/56/68 ms

Sending 5, 100-byte ICMP Echos to 200.110.12.1, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 44/57/68 ms

Sending 5, 100-byte ICMP Echos to 200.110.12.129, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 52/63/80 ms

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 223.255.255.1, timeout is 2 seconds:

Packet sent with a source address of 128.45.0.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 4/34/64 ms

## 5. ISP politika

### 5.1 Primárne linky R3 – R8 a R4 – R10

Úlohou bolo zabezpečiť, aby bola pre smerovače R8 a R9 zvolená primárna trasa z AS 65001 cez smerovač R3, a nie cez R4. Aby sme zabezpečili takéto správanie, rozhodli sme sa zvýhodniť linku medzi R3 a R8 pomocou hodnoty Local Preference. Na smerovači R8 bolo potrebné vytvoriť route-mapu, v ktorej sme nastavili hodnotu local preference na 150 (default hodnota je 100 a v tomto prípade platí, že berie vyššia). Následne bolo potrebné aplikovať route-mapu na suseda, v našom prípade na R3.

**3R8:** route-map R8-out permit 10

**set local-preference 150**

*aplikovanie route-mapy:* neighbor 200.110.255.241 route-map R8-out in

Aby bola úloha splnená, bolo potrebné ešte nastaviť zvýhodnenie linky v opačnom smere, z R3 na R8. Na tento prípad sme využili komunity. Smerovače R8 a R9 označili všetku prevádzku smerovanú z vnútra AS 65001.

**3R8:**

access-list 100 permit ip any any

route-map R8-R3 permit 10

match ip address 100

**set community 65001:2001 additive**

neighbor 200.110.255.241 route-map R8-R3 out

neighbor 200.110.255.241 send-community

**3R9:**

access-list 100 permit ip any any

route-map R9-R4 permit 10

match ip address 100

**set community 65001:2001 additive**

neighbor 200.110.255.237 route-map R9-R4 out

neighbor 200.110.255.237 send-community

Takto označenú prevádzku sme museli zachytávať na smerovači R3, ktorý mal byť hlavnou výstupnou cestou z AS 65001. Vytvorili sme route-mapu, v ktorej sme nastavili rovnako hodnotu local preference na 150.

**3R3:**

ip community-list 2 permit 65001:2001

route-map R3-R8 permit 10

match community 2

**set local-preference 150**

neighbor 200.110.255.242 route-map R3-R8 in

Aby sa komunity zobrazovali v prijateľnom tvare, tak ako ho zadávame (65001:2001), bolo potrebné na všetkých smerovačoch zadať príkaz *ip bgp-community new-format*. Aby sa prejavili nakonfigurované zmeny, bolo potrebné resetovať BGP príkazom *clear ip bgp \* in/out*.

Správne nakonfigurovanie overíme zobrazením BGP tabuľky pre sieť 200.110.12.128 zo smerovača R3, kde by sme mali vidieť nastavenú hodnotu local preference na 150 a takisto aj komunitu 65001:2001. Rovnaký princíp využijeme aj pri kontrole na R9, kde by mala byť nastavená primárna cesta z AS 65001 cez smerovač R8. Príkazom traceroute z R9 na smerovač R5 potvrdíme správnosť konfigurácie.

```
3R3#sh ip bgp 200.110.12.128
```

```
BGP routing table entry for 200.110.12.0/24, version 42
Paths: (1 available, best #1, table Default-IP-Routing-Table)
Flag: 0x820
Advertised to update-groups:
  1
  65001, (aggregated by 65001 10.255.255.8)
    200.110.255.242 from 200.110.255.242 (10.255.255.8)
      Origin IGP, metric 0, localpref 150, valid, external, atomic-aggregate, best
      Community: 65001:2001
```

```
3R9#sh ip bgp 200.110.0.0/21
```

```
BGP routing table entry for 200.110.0.0/21, version 42
Paths: (2 available, best #1, table Default-IP-Routing-Table)
Advertised to update-groups:
  1
  110, (aggregated by 110 200.110.2.1)
    200.110.255.241 (metric 10) from 10.255.255.8 (200.110.12.1)
      Origin IGP, metric 0, localpref 150, valid, internal, atomic-aggregate, best
  110, (aggregated by 110 200.110.4.1)
    200.110.255.237 from 200.110.255.237 (200.110.4.1)
      Origin IGP, metric 0, localpref 100, valid, external, atomic-aggregate
```

```
3R9#traceroute 128.45.0.1 source 200.110.12.129
```

```
Type escape sequence to abort.
Tracing the route to 128.45.0.1

 1 10.110.89.8 32 msec 20 msec 16 msec
 2 200.110.255.241 88 msec 52 msec 40 msec
 3 10.110.23.2 56 msec 80 msec 104 msec
 4 200.110.255.254 92 msec 92 msec 128 msec
```

Posledným krokom v tomto bode bolo nastaviť primárnu linku z R10. Zo zadania vyplýva, že primárnou linkou, ktorou majú byť smerované údaje má byť spojenie R4-R10. Zo strany R10 zvýhodníme linku pomocou hodnoty local preference, ktorú nastavíme na 180.

```
3R10: route-map R10-R4 permit 10
      set local-preference 180
```

Týmto krokom sme zabezpečili požadované správanie smerom z R10. Následne bolo treba ošetriť aj opačný smer, z R4 na R10. V tomto bode sme využili možnosť znevýhodniť linku R7-R10 pomocou AS\_PATH. Využitím route-mapy sme si označili všetku prevádzku, ktorá je vytváraná v AS 5005 a pomocou príkazu *prepend* sme umelo natiahli cestu z AS 5005.

### 3R10:

```
ip as-path access-list 10 permit ^$ (všetko čo vychádza z daného AS)
```

```
route-map AS_PATH permit 10
```

```
match as-path 10
```

```
set as-path prepend 5005 5005 5005
```

```
neighbor 200.33.255.245 route-map AS_PATH out
```

Týmto pádom, by smerovač R7 mal vidieť cestu do R10 cez priamu linku ako menej výhodnú, keďže prechádza cez viacero AS. Po resetovaní BGP by sa mali prejaviť zmeny, a je nutné overiť ich korektnosť.

### 3R10#sh ip bgp 128.45.0.0

BGP routing table entry for 128.45.0.0/24, version 27

Paths: (2 available, best #1, table Default-IP-Routing-Table)

Not advertised to any peer

110 4502

200.110.255.245 from 200.110.255.245 (200.110.4.1)

Origin IGP, localpref 180, valid, external, best

330 4502

200.33.255.245 from 200.33.255.245 (200.33.7.1)

Origin IGP, localpref 100, valid, external

### 3R7#sh ip bgp

| Network           | Next Hop       | Metric | LocPrf | Weight | Path                  |
|-------------------|----------------|--------|--------|--------|-----------------------|
| *>i64.34.0.0/24   | 200.33.255.253 | 0      | 100    | 0      | 3401 i                |
| *>i128.45.0.0/24  | 200.33.255.249 | 0      | 100    | 0      | 4502 i                |
| r>i200.33.6.0     | 10.255.255.6   | 0      | 100    | 0      | i                     |
| *> 200.33.7.0     | 0.0.0.0        | 0      |        | 32768  | i                     |
| *>i200.110.0.0/21 | 200.33.255.253 | 0      | 100    | 0      | 3401 110 i            |
| *>i200.110.12.0   | 200.33.255.249 | 0      | 100    | 0      | 4502 110 i            |
| *>i223.255.255.0  | 200.33.255.253 | 0      | 100    | 0      | 3401 110 5005 i       |
| *                 | 200.33.255.246 | 0      |        | 0      | 5005 5005 5005 5005 i |

```
3R7#sh ip bgp 223.255.255.0
```

BGP routing table entry for 223.255.255.0/24, version 29  
Paths: (2 available, best #1, table Default-IP-Routing-Table)

Advertised to update-groups:

1

3401 110 5005

200.33.255.253 (metric 10) from 10.255.255.6 (200.33.6.1)

Origin IGP, metric 0, localpref 100, valid, internal, best

5005 5005 5005 5005

200.33.255.246 from 200.33.255.246 (223.255.255.1)

Origin IGP, metric 0, localpref 100, valid, external

## 5.2 AS 5005 nesmie byť nikdy transit

Pre úspešné zvládnutie tejto úlohy bolo potrebné na R10 nakonfigurovať access-list, ktorý povoľoval všetku prevádzku vychádzajúcu z R10 (AS 5005). Následne ho zahrnúť do route-mapy a tú aplikovať na oboch susedov, R4 aj R7.

**3R10:**

```
ip as-path access-list 1 permit ^$
```

```
route-map NO_TRANSIT permit 10  
match as-path 1
```

```
neighbor 200.33.255.245 route-map NO_TRANSIT out  
neighbor 200.110.255.245 route-map NO_TRANSIT out
```

Obdobným štýlom nakonfigurujeme aj smerovače R4 a R7, len s tou výnimkou, že v access-liste povolíme iba prevádzku, ktorá začína v R10.

**3R4:**

```
ip as-path access-list 1 permit _5005$
```

```
route-map NO_TRANSIT permit 10  
match as-path 1
```

```
neighbor 200.110.255.246 route-map NO_TRANSIT in
```

**3R7:**

```
ip as-path access-list 1 permit _5005$
```

```
route-map NO_TRANSIT permit 10  
match as-path 1
```

```
neighbor 200.33.255.246 route-map NO_TRANSIT in
```

Po resetovaní by sa mali prejaviť želané zmeny, a overenie vykonáme pomocou príkazu traceroute z R4 na R7. Overíme aj BGP tabuľky, či najlepšie cesty R4-R7 neprechádzajú cez R10.

**3R7#sh ip bgp 200.110.0.0/21**

Advertised to update-groups:

1

**3401 110**, (aggregated by 110 200.110.2.1)

200.33.255.253 (metric 10) from 10.255.255.6 (200.33.6.1)

Origin IGP, metric 0, localpref 100, valid, internal, atomic-aggregate, **best**

**3R4#sh ip bgp 200.33.7.0/24**

Advertised to update-groups:

1 3

**4502 330**

200.110.255.254 (metric 10) from 10.255.255.2 (200.110.2.1)

Origin IGP, metric 0, localpref 100, valid, internal, **best**

**3R4#traceroute 200.33.7.1 source 200.110.4.1**

1 **10.110.24.2** 20 msec 24 msec 36 msec

2 200.110.255.254 24 msec 44 msec 0 msec

3 200.33.255.250 64 msec 76 msec 44 msec

4 10.110.67.7 108 msec \* 20 msec

### 5.3 Peering iba pre ISP1 a ISP2

V záujme lokálnych ISP bolo potrebné nastaviť, aby prevádzka nebola smerovaná z jedného ISP do Upstream ISP cez druhého providera. Keďže sú tieto linky spoplatnené, chceme zabezpečiť, aby sa nemohli zneužiť. Na smerovači R1, ktorý reprezentuje Upstream ISP, budeme označovať jeho vlastnú prevádzku.

**3R1:**

```
route-map R1-COM permit 10
  set community 3401:1001 additive
```

```
neighbor 200.33.255.254 send-community
neighbor 200.33.255.254 route-map R1-COM out
neighbor 200.110.255.250 send-community
neighbor 200.110.255.250 route-map R1-COM out
```

Následne bolo nutné túto značkovанú prevádzku zachytávať. Na R2 a R6 sme vytvorili community-list, a následne sme v route-mape zakázali šírenie tejto prevádzky smerom na R5.

**3R2:**

```
ip community-list 2 permit 3401:1001
```

```
route-map COMUN_UPS deny 10
  match community 2
```

```
neighbor 200.110.255.254 route-map COMUN_UPS out
```

**3R6:**

```
ip community-list 2 permit 3401:1001

route-map COMUN_UPS deny 10
  match community 2

neighbor 200.33.255.249 route-map COMUN_UPS out
```

Týmto krokom sme zabezpečili, aby sa ISP1 zo smerovača R2 nedostal na R1 cez smerovač R6, a opačne. Konfiguráciu overíme na smerovači R5, ktorý ukazuje len jedinú cestu do R1, a to jeho priamu linku.

**3R5#sh ip bgp 64.34.0.0/24**

```
BGP routing table entry for 64.34.0.0/24, version 11
Paths: (1 available, best #1, table Default-IP-Routing-Table)
Advertised to update-groups:
  1

3401
  64.34.255.253 from 64.34.255.253 (64.34.0.1)
    Origin IGP, metric 0, localpref 100, valid, external, best
```

## 5.4 Distribúovať iba default, AS 5005 a peering prefixy do AS 65001

Na smerovači R2 sme smerom nadol (na R3,R4) distribuovali len default route, čím sme zjednodušili BGP tabuľky a odbremenili sme ich od prefixov posielaných z Upstream ISP.

**3R2:**

```
neighbor 10.255.255.3 default-originate
neighbor 10.255.255.4 default-originate
```

Overenie:

**3R8#sh ip bgp**

|     | Network         | Next Hop        | Metric | LocPrf | Weight | Path       |
|-----|-----------------|-----------------|--------|--------|--------|------------|
| *>  | 0.0.0.0         | 200.110.255.241 |        | 150    | 0      | 110 i      |
| *>  | 200.110.0.0/21  | 200.110.255.241 |        | 150    | 0      | 110 i      |
| *>  | 200.110.3.0     | 200.110.255.241 | 0      | 150    | 0      | 110 i      |
| s>  | 200.110.12.0/25 | 0.0.0.0         | 0      |        | 32768  | i          |
| *>  | 200.110.12.0    | 0.0.0.0         |        |        | 32768  | i          |
| * i |                 | 10.255.255.9    | 0      | 100    | 0      | i          |
| *>  | 223.255.255.0   | 200.110.255.241 |        | 150    | 0      | 110 5005 i |

Distribúciu prefixov z AS 5005 a Peering centra sme vyriešili pomocou komunit. Na smerovačoch R5 a R10 sme začali značkovať prevádzku a šírili ju vďaka príkazu *send-community* na smerovačoch, ktoré sa nachádzali na trase do AS 65001.

#### 3R10:

```
route-map R10-COM permit 10
  set community 5005:5001 additive

neighbor 200.33.255.245 send-community
neighbor 200.33.255.245 route-map R10-COM out
```

#### 3R5:

```
route-map R5-COM permit 10
  set community 4502:6001 additive

neighbor 200.110.255.253 send-community
neighbor 200.110.255.253 route-map R5-COM out
```

Po tomto kroku by malo byť všetko potrebné nastavené. Overenie vykonáme opäť prezretím BGP tabuľky, kde sa už nebudú nachádzať ani záznamy z AS 5005 (223.255.255.0). Zabezpečená je aj plná konektivita medzi smerovačmi R8 a R5,R10.

#### 3R8#sh ip bgp

|     | Network         | Next Hop        | Metric | LocPrf | Weight | Path  |
|-----|-----------------|-----------------|--------|--------|--------|-------|
| *>  | 0.0.0.0         | 200.110.255.241 | 150    | 0      | 0      | 110 i |
| *>  | 200.110.0.0/21  | 200.110.255.241 |        | 150    | 0      | 110 i |
| *>  | 200.110.3.0     | 200.110.255.241 | 0      | 150    | 0      | 110 i |
| s>  | 200.110.12.0/25 | 0.0.0.0         | 0      |        | 32768  | i     |
| *>  | 200.110.12.0    | 0.0.0.0         |        |        | 32768  | i     |
| * i |                 | 10.255.255.9    | 0      | 100    | 0      | i     |

#### 3R8#traceroute 128.45.0.1 source 200.110.12.1

```
1 200.110.255.241 [AS 110] 16 msec 48 msec 28 msec
2 10.110.23.2 [AS 110] 64 msec 64 msec 60 msec
3 200.110.255.254 [AS 110] 84 msec 104 msec 80 msec
```

#### 3R8#traceroute 223.255.255.1 source 200.110.12.1

```
1 200.110.255.241 [AS 110] 16 msec 32 msec 20 msec
2 10.110.34.4 [AS 110] 72 msec 60 msec 32 msec
3 200.110.255.246 [AS 110] 80 msec * 64 msec
```

## 5.5 Overiť, či je možné odkloniť celú prevádzku na linke R4-R10 v prípade údržby

V prípade výpadku linky medzi R4 a R10 by mala byť zabezpečená plná konektivita medzi všetkými smerovačmi. Overenie môžeme spraviť pomocou *pingu* na všetky šírené prefixy, a prípadne príkazom *traceroute* z R8 na R10. Za normálnych okolností je prevádzka z R8 smerovaná cez R3 → R4 → R10. Pri výpadku linky, by mala byť nájdená záložná cesta cez Peeringové centrum (R3 → R2 → R5 → R6 → R7 → R10).



```

3R8(tcl)#foreach address {
+>64.34.0.1
+>200.110.2.1
+>200.110.3.1
+>200.110.4.1
+>128.45.0.1
+>200.33.6.1
+>200.33.7.1
+>200.110.12.129
+>223.255.255.1
+>} {
+>ping $address source 200.110.12.1 }

```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 64.34.0.1, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 84/94/100 ms

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 200.110.2.1, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 28/54/84 ms

Sending 5, 100-byte ICMP Echos to 200.110.3.1, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 12/26/48 ms

Sending 5, 100-byte ICMP Echos to 200.110.4.1, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 40/59/88 ms

Sending 5, 100-byte ICMP Echos to 128.45.0.1, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 64/91/112 ms

Sending 5, 100-byte ICMP Echos to 200.33.6.1, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 68/90/116 ms

Sending 5, 100-byte ICMP Echos to 200.33.7.1, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 108/121/136 ms

Sending 5, 100-byte ICMP Echos to 200.110.12.129, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 8/24/44 ms

Sending 5, 100-byte ICMP Echos to 223.255.255.1, timeout is 2 seconds:

Packet sent with a source address of 200.110.12.1

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 96/119/140 ms

**3R8#traceroute 223.255.255.1 source 200.110.12.1**

```
1 200.110.255.241 [AS 110] 8 msec 56 msec 24 msec
2 10.110.23.2 [AS 110] 40 msec 72 msec 72 msec
3 200.110.255.254 [AS 110] 92 msec 88 msec 100 msec
4 200.33.255.250 [AS 110] 88 msec 60 msec 120 msec
5 10.110.67.7 [AS 110] 104 msec 144 msec 88 msec
6 200.33.255.246 [AS 110] 124 msec * 104 msec
```

## 5.6 Overiť funkčnosť nastavenia politiky vhodnými výpadkami liniek

Správne nastavenie politiky sme demonštrovali v predošlej úlohe, keď nastal výpadok na linke R4-R10. Tento výpadok nemal nijaký vplyv na konektivitu v celej sieti, keďže bola nájdená záložná trasa. Otestovať sme sa rozhodli ešte výpadky liniek medzi smerovačmi R2 a R5, respektíve medzi R6 a R5. Z výpisu je vidieť, že siete sú plne dostupné aj pri výpadku.

**3R2#traceroute 223.255.255.1 source 200.110.2.1**

```
1 200.110.255.249 20 msec 28 msec 44 msec
2 200.33.255.254 36 msec 24 msec 28 msec
3 10.110.67.7 40 msec 52 msec 36 msec
4 200.33.255.246 52 msec 60 msec *
```

**3R6#traceroute 200.110.12.1 source 200.33.6.1**

```
1 200.33.255.253 0 msec 16 msec 12 msec
2 200.110.255.250 24 msec 32 msec 8 msec
3 10.110.23.3 80 msec 64 msec 60 msec
4 200.110.255.242 88 msec * 64 msec
```